

```
In [111... import numpy as np

import pandas as pd #excellent for dataset manipulation

# for data visulization
import matplotlib.pyplot as plt

#stats visualization
import seaborn as sns

#Labelencoding to convert categorical data into Lowlevel Language
from sklearn.preprocessing import LabelEncoder

#scaling data
from sklearn.preprocessing import StandardScaler

#data partions
from sklearn.model_selection import train_test_split

#algorithms
from sklearn.linear_model import LogisticRegression

from sklearn.tree import DecisionTreeClassifier

from sklearn.ensemble import RandomForestClassifier

from xgboost import XGBClassifier

#accuracy confusion matric and classification report
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

import warnings

# To ignore all warnings
warnings.filterwarnings("ignore")
```

```
In [112... df = pd.read_csv(r"C:\Users\lenovo\Desktop\NIT FILES\New folder\diabetes_predict
```

```
In [113... df.head()
```

```
Out[113... 
```

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blo
0	Female	80.0	0	1	never	25.19	6.6	
1	Female	54.0	0	0	No Info	27.32	6.6	
2	Male	28.0	0	0	never	27.32	5.7	
3	Female	36.0	0	0	current	23.45	5.0	
4	Male	76.0	1	1	current	20.14	4.8	

```
In [114... df.corr(numeric_only=True)
```

Out[114...

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level
age	1.000000	0.251171	0.233354	0.337396	0.101354	
hypertension	0.251171	1.000000	0.121262	0.147666	0.080939	
heart_disease	0.233354	0.121262	1.000000	0.061198	0.067589	
bmi	0.337396	0.147666	0.061198	1.000000	0.082997	
HbA1c_level	0.101354	0.080939	0.067589	0.082997	1.000000	
blood_glucose_level	0.110672	0.084429	0.070066	0.091261	0.166733	
diabetes	0.258008	0.197823	0.171727	0.214357	0.400660	



In [115...

df.shape

Out[115...

(100000, 9)

In [116...

```
for column in df.columns:
    unique_values = df[column].unique()
    print('Column "{}" has unique values: {}'.format(column, unique_values))
```

Column "gender" has unique values: ['Female' 'Male' 'Other']

Column "age" has unique values: [80. 54. 28. 36. 76. 20. 44. 79. 42. 32. 53. 78.

```
67. 15. 37. 40. 5. 69. 72. 4. 30. 45. 43. 50.
41. 26. 34. 73. 77. 66. 29. 60. 38. 3. 57. 74.
19. 46. 21. 59. 27. 13. 56. 2. 7. 11. 6. 55.
9. 62. 47. 12. 68. 75. 22. 58. 18. 24. 17. 25.
0.08 33. 16. 61. 31. 8. 49. 39. 65. 14. 70. 0.56
48. 51. 71. 0.88 64. 63. 52. 0.16 10. 35. 23. 0.64
1.16 1.64 0.72 1.88 1.32 0.8 1.24 1. 1.8 0.48 1.56 1.08
0.24 1.4 0.4 0.32 1.72 1.48]
```

Column "hypertension" has unique values: [0 1]

Column "heart_disease" has unique values: [1 0]

Column "smoking_history" has unique values: ['never' 'No Info' 'current' 'former' 'ever' 'not current']

Column "bmi" has unique values: [25.19 27.32 23.45 ... 59.42 44.39 60.52]

Column "HbA1c_level" has unique values: [6.6 5.7 5. 4.8 6.5 6.1 6. 5.8 3.5 6.2 4. 4.5 9. 7. 8.8 8.2 7.5 6.8]

Column "blood_glucose_level" has unique values: [140 80 158 155 85 200 145 100 130 160 126 159 90 260 220 300 280 240]

Column "diabetes" has unique values: [0 1]

In [117...

df["smoking_history"].value_counts()

Out[117...

```
smoking_history
No Info      35816
never        35095
former        9352
current       9286
not current   6447
ever          4004
Name: count, dtype: int64
```

In [118...

df["smoking_history"].value_counts()/len(df)

```
Out[118...] smoking_history
No Info      0.35816
never        0.35095
former       0.09352
current      0.09286
not current  0.06447
ever         0.04004
Name: count, dtype: float64
```

```
In [119...] # Replaceing No Info columns with pd.NA
df['smoking_history'] = df['smoking_history'].replace('No Info', pd.NA)

# Replace missing values with the mode it is string so we are using mode
mode_value = df['smoking_history'].mode()[0]
df['smoking_history'] = df['smoking_history'].fillna(mode_value) #filling no inf

# Printing the updated value counts
print(df['smoking_history'].value_counts())
```

```
smoking_history
never          70911
former         9352
current        9286
not current    6447
ever           4004
Name: count, dtype: int64
```

```
In [120...] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                100000 non-null  object
1   age                   100000 non-null  float64
2   hypertension          100000 non-null  int64
3   heart_disease         100000 non-null  int64
4   smoking_history       100000 non-null  object
5   bmi                   100000 non-null  float64
6   HbA1c_level           100000 non-null  float64
7   blood_glucose_level   100000 non-null  int64
8   diabetes              100000 non-null  int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB
```

```
In [121...] df.describe()
```

Out[121...

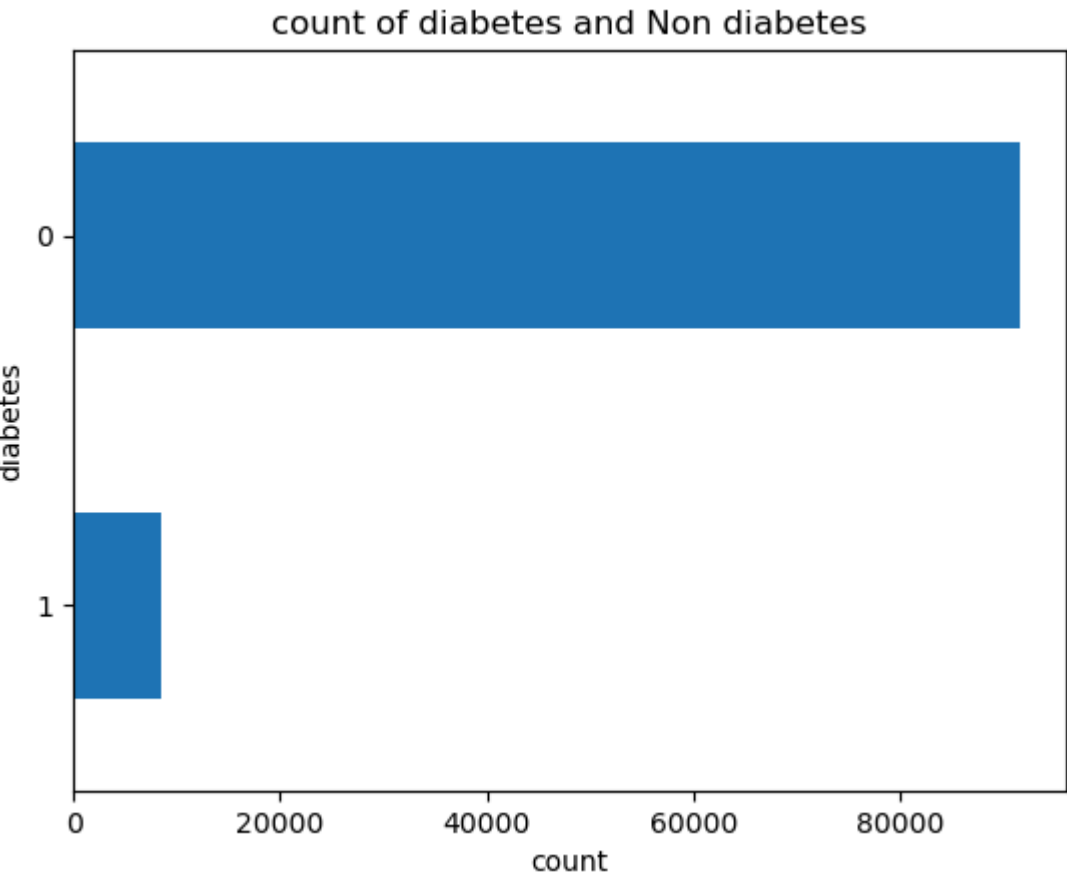
	age	hypertension	heart_disease	bmi	HbA1c_level	blood_sugar
count	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000
mean	41.885856	0.07485	0.039420	27.320767	5.527507	120.917188
std	22.516840	0.26315	0.194593	6.636783	1.070672	54.156491
min	0.080000	0.00000	0.000000	10.010000	3.500000	62.000000
25%	24.000000	0.00000	0.000000	23.630000	4.800000	99.000000
50%	43.000000	0.00000	0.000000	27.320000	5.800000	119.000000
75%	60.000000	0.00000	0.000000	29.580000	6.200000	131.000000
max	80.000000	1.00000	1.000000	95.690000	9.000000	199.000000

In [122...

```
df["bmi"] = [float(str(i).replace(",","")) for i in df["bmi"]]
```

In [123...

```
df['diabetes'].value_counts().plot(kind='barh')
plt.xlabel('count')
plt.ylabel('diabetes')
plt.title('count of diabetes and Non diabetes')
plt.gca().invert_yaxis()
plt.show()
```



In [124...

```
df['diabetes'].value_counts()/len(df)
```

```
Out[124...] diabetes
0    0.915
1    0.085
Name: count, dtype: float64
```

```
In [125...] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                100000 non-null object
1   age                   100000 non-null float64
2   hypertension          100000 non-null int64
3   heart_disease         100000 non-null int64
4   smoking_history       100000 non-null object
5   bmi                   100000 non-null float64
6   HbA1c_level           100000 non-null float64
7   blood_glucose_level   100000 non-null int64
8   diabetes              100000 non-null int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB
```

```
In [126...] le = LabelEncoder()
le
```

```
Out[126...] ▼ LabelEncoder ⓘ ?
LabelEncoder()
```


```
In [127...] Label_encod_columns=['gender', 'smoking_history']
df[Label_encod_columns]= df[Label_encod_columns].apply(le.fit_transform)
```

```
In [128...] df.head(3)
```

```
Out[128...]
   gender  age  hypertension  heart_disease  smoking_history  bmi  HbA1c_level  blood_glucose_level
0      0  80.0             0             1             3  25.19             6.6
1      0  54.0             0             0             3  27.32             6.6
2      1  28.0             0             0             3  27.32             5.7
```


```
In [129...] sns.boxplot(data=df[['age', 'blood_glucose_level', 'bmi']])
```

```
Out[129...] <Axes: >
```




```
In [130...] sns.boxplot(data=df['HbA1c_level'])
```

```
Out[130...] <Axes: ylabel='HbA1c_level'>
```

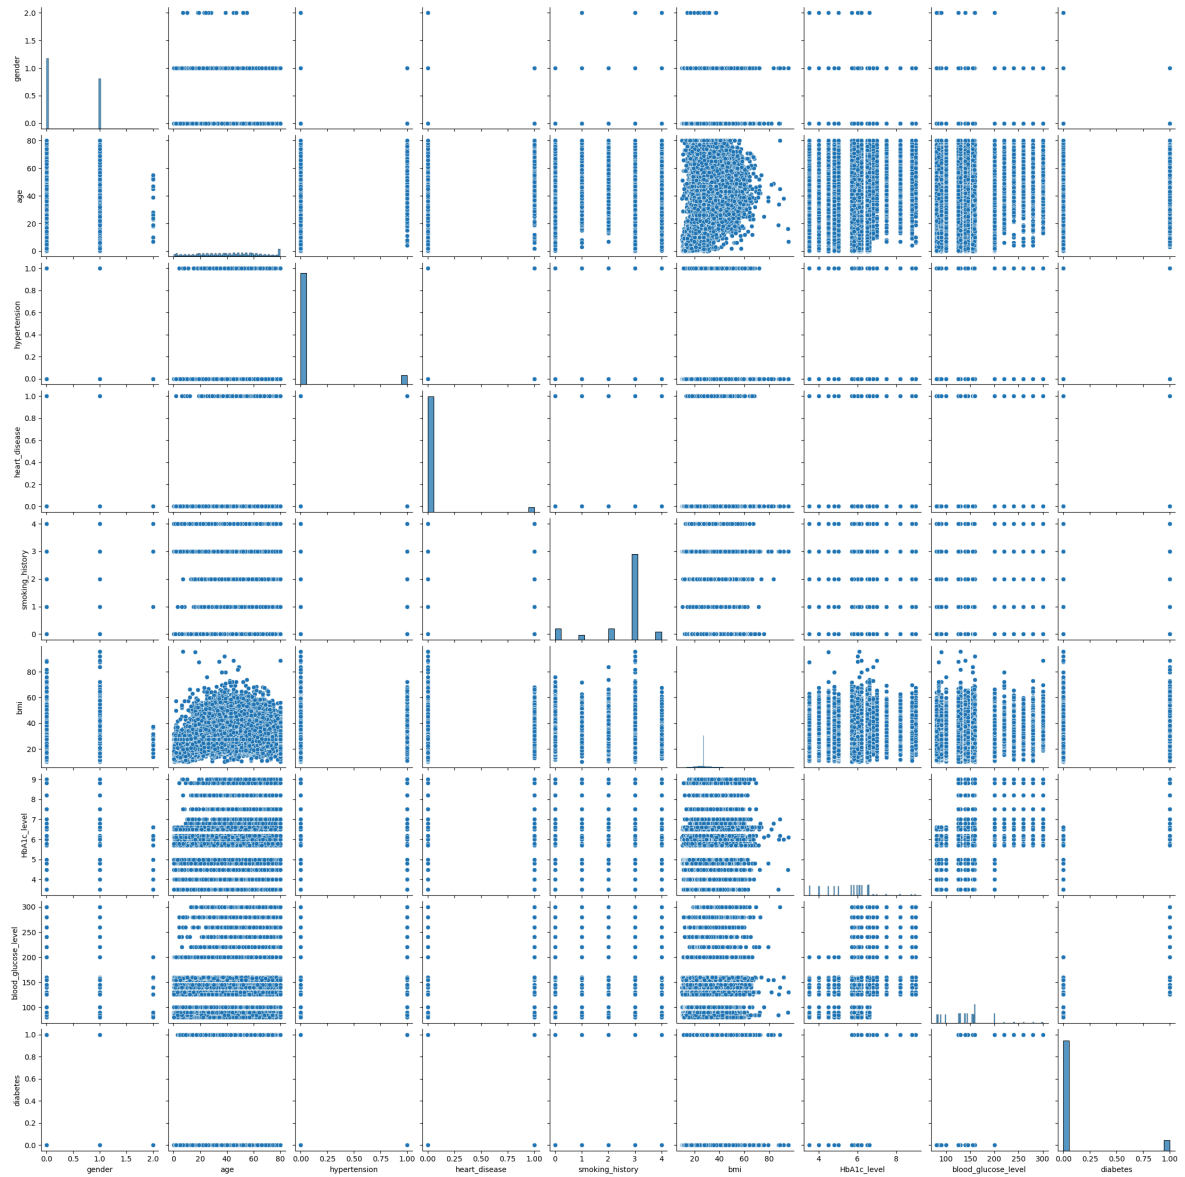


```
In [131...] sns.lmplot(data=df, x='blood_glucose_level', y='diabetes', fit_reg=False)
```

```
Out[131...] <seaborn.axisgrid.FacetGrid at 0x1f28c4e6330>
```



```
In [132... sns.pairplot(df)  
plt.show()
```



In [133...]

```
'''when age increase hypertension and heart disease ,blood_glucose_level and diabetes
relationship between them

*bmi

*HbA1c_level

*blood_glucose_level

these four parameters have relationship between each other

*gender and smoking history it does not effect on diabetes
'''
```

Out[133...]

```
'when age increase hypertension and heart disease ,blood_glucose_level and diabetes
and age and also there is a \n relationship between them\n\n *bmi\n\n *HbA1c_level\n\n *blood_glucose_level\n\n these four parameters
have relationship between each other\n\n *gender and smoking history
it does not effect on diabetes\n'
```

In [134...]

```
df.corr()
```



Out[134...

	gender	age	hypertension	heart_disease	smoking_history
gender	1.000000	-0.030656	0.014203	0.077696	-0.044081
age	-0.030656	1.000000	0.251171	0.233354	-0.098969
hypertension	0.014203	0.251171	1.000000	0.121262	-0.048631
heart_disease	0.077696	0.233354	0.121262	1.000000	-0.048253
smoking_history	-0.044081	-0.098969	-0.048631	-0.048253	1.000000
bmi	-0.022994	0.337396	0.147666	0.061198	-0.087735
HbA1c_level	0.019957	0.101354	0.080939	0.067589	-0.017534
blood_glucose_level	0.017199	0.110672	0.084429	0.070066	-0.022985
diabetes	0.037411	0.258008	0.197823	0.171727	-0.049841

In [135...

```
plt.figure(figsize=(20,8))
df.corr()['diabetes'].sort_values(ascending=False).plot(kind='bar')
```

Out[135...



In [136...

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                100000 non-null  int32
1   age                   100000 non-null  float64
2   hypertension          100000 non-null  int64
3   heart_disease         100000 non-null  int64
4   smoking_history       100000 non-null  int32
5   bmi                   100000 non-null  float64
6   HbA1c_level           100000 non-null  float64
7   blood_glucose_level   100000 non-null  int64
8   diabetes              100000 non-null  int64
dtypes: float64(3), int32(2), int64(4)
memory usage: 6.1 MB
```

```
In [137... X = df.loc[:, 'age':'heart_disease'].join(df.loc[:, 'bmi':'blood_glucose_level'])
X
```

```
Out[137...      age  hypertension  heart_disease  bmi  HbA1c_level  blood_glucose_level
0  80.0              0              1  25.19           6.6           140
1  54.0              0              0  27.32           6.6            80
2  28.0              0              0  27.32           5.7           158
3  36.0              0              0  23.45           5.0           155
4  76.0              1              1  20.14           4.8           155
...   ...              ...              ...   ...           ...           ...
99995  80.0              0              0  27.32           6.2            90
99996   2.0              0              0  17.37           6.5           100
99997  66.0              0              0  27.83           5.7           155
99998  24.0              0              0  35.42           4.0           100
99999  57.0              0              0  22.43           6.6            90
```

100000 rows × 6 columns

```
In [138... Y = df.loc[:, 'diabetes']
Y
```

```
Out[138... 0      0
1      0
2      0
3      0
4      0
..
99995  0
99996  0
99997  0
99998  0
99999  0
Name: diabetes, Length: 100000, dtype: int64
```

Data Partition

```
In [139... # splitting training and testing data in 70 30 ratio testing size is 0.3 random_s
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.2,random_state=0)
```

```
In [140... X_train.head() #printing X_train data
```

```
Out[140...
      age  hypertension  heart_disease  bmi  HbA1c_level  blood_glucose_level
10382   2.0             0             0  16.45           6.2             159
73171  55.0             0             0  24.59           6.0             130
30938  24.0             0             0  21.77           4.5             130
99310  30.0             0             0  27.32           6.2             159
58959  13.0             0             0  18.37           6.5             130
```

```
In [141... print('Shape of Train data')
print(X_train.shape)
print(Y_train.shape)
print('Shape of Testing data')
print(X_test.shape)
print(Y_test.shape)
```

Shape of Train data
(80000, 6)
(80000,)
Shape of Testing data
(20000, 6)
(20000,)

```
In [142... ss=StandardScaler() #activating StandardScaler()
ss
```

```
Out[142...
▼ StandardScaler ⓘ ?
StandardScaler()
```

```
In [143... X_train_scaled=ss.fit_transform(X_train) #scaling X_train data
```

```
In [144... if len(X_test.shape) == 1: #if x is 1d array
    X_test = X_test.values.reshape(-1, 1) #converting to 2d array
X_test_scaled = ss.fit_transform(X_test) #scaling X_test data
```

```
In [145... model_lr=LogisticRegression() #activating Logistic Regression
```

```
In [146... model_lr.fit(X_train_scaled,Y_train) #training logistic regression model
```

```
Out[146... LogisticRegression
```

```
LogisticRegression()
```

```
In [147... y_pred=model_lr.predict(X_test_scaled) #predecting y_test data
y_pred[:10]
```

```
Out[147... array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int64)
```

```
In [148... Y_test[:10] # actual Y_test data
```

```
Out[148... 3582    0
60498    0
53227    0
21333    0
3885     0
51521    0
84261    0
10685    1
59948    0
41032    0
Name: diabetes, dtype: int64
```

```
In [149... accuracy_score(y_pred,Y_test) #accuracy_score
```

```
Out[149... 0.95975
```

```
In [150... print(classification_report(y_pred,Y_test)) #classifiaction_report
```

	precision	recall	f1-score	support
0	0.99	0.97	0.98	18736
1	0.63	0.86	0.73	1264
accuracy			0.96	20000
macro avg	0.81	0.91	0.85	20000
weighted avg	0.97	0.96	0.96	20000

```
In [151... Y_train.value_counts() #data is highly imblancing
```

```
Out[151... diabetes
0    73203
1     6797
Name: count, dtype: int64
```

```
In [152... value_counts=Y_train.value_counts()


plt.figure(figsize=(16, 8))

plt.pie(value_counts, labels=value_counts.index, autopct='%1.2f%%', startangle=1

plt.title('Distribution of Y_train')


plt.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.
```

```
plt.show()
```



```
In [153... from imblearn.over_sampling import SMOTE # using smote function to balance our s
smote=SMOTE()
X_ovs,Y_ovs=smote.fit_resample(X,Y) #passing X and y variables to it to balance
fig, oversp = plt.subplots()
oversp.pie( Y_ovs.value_counts(), autopct='%.2f')
oversp.set_title("Over-sampling")
plt.show()
```

Over-sampling



```
In [154... # Dividing our resampling data into 70 30 ratio

Xr_train,Xr_test,Yr_train,Yr_test=train_test_split(X_ovs,Y_ovs,train_size=0.7,ra
```

```
In [155... print('train data shape')

print(Xr_train.shape)

print(Yr_train.shape)

print('test data shape')

print(Xr_test.shape)

print(Yr_test.shape)
```

```
train data shape
(128099, 6)
(128099,)
test data shape
(54901, 6)
(54901,)
```

```
In [156... print('Y_train and Y_test value_count')
print(Yr_train.value_counts())
print(Yr_test.value_counts())
```

```
Y_train and Y_test value_count
diabetes
0    64131
1    63968
Name: count, dtype: int64
diabetes
1    27532
0    27369
Name: count, dtype: int64
```

```
In [157... ss=StandardScaler()

ss
```

```
Out[157... ▼ StandardScaler ⓘ ?
StandardScaler()
```

```
In [158... data=Xr_train,Xr_test

Xr_train_sc=ss.fit_transform(Xr_train) # scaling our resampling data xr train

Xr_test_sc=ss.fit_transform(Xr_test) # scaling our resamplig xr_test data
```

```
In [159... Xr_train_scaled = pd.DataFrame(Xr_train_sc) #Xr_train_scaled converting into the

print(Xr_train_scaled.shape)
```

```
Xr_train_scaled.head()
print(Yr_train.shape)
```

```
(128099, 6)
(128099,)
```

```
In [160... Xr_test_scaled=pd.DataFrame(Xr_test_sc) #Xr_test converting into the dataframe

print(Xr_test_scaled.shape)
Xr_test_scaled.head()
```

```
(54901, 6)
```

```
Out[160...      0      1      2      3      4      5
0 -1.233879 -0.293163 -0.204896 -0.443557 -0.256457 -0.325757
1 -1.093560 -0.293163 -0.204896 -0.406660  0.370289 -0.063854
2 -1.465145 -0.293163 -0.204896 -0.290170  0.370289 -1.460670
3 -0.768423  3.411077 -0.204896  0.284151  0.370289 -1.373369
4 -1.372249 -0.293163 -0.204896 -0.290170 -2.154733 -1.111466
```

```
In [161... model_lk=LogisticRegression()

model_lk.fit(Xr_train_scaled,Yr_train) #training the model
```

```
Out[161... LogisticRegression
LogisticRegression()
```

```
In [162... y_pred_lr=model_lk.predict(Xr_test_scaled) #predicting yr_test data
y_pred_lr[:10]
```

```
Out[162... array([0, 0, 0, 0, 0, 1, 0, 0, 0, 1], dtype=int64)
```

```
In [163... Yr_test[:10]
```

```
Out[163... 180328    1
573        0
13494       0
93981       0
75389       0
180973     1
71021       0
19293       0
16393       0
121419     1
Name: diabetes, dtype: int64
```

```
In [164... #classification_report for predict value and original value

print(classification_report(y_pred_lr,Yr_test))
```

	precision	recall	f1-score	support
0	0.88	0.88	0.88	27489
1	0.88	0.88	0.88	27412
accuracy			0.88	54901
macro avg	0.88	0.88	0.88	54901
weighted avg	0.88	0.88	0.88	54901

```
In [165... #confusion_matrix for predict value and original value

confusion_matrix(y_pred_lr,Yr_test)
```

```
Out[165... array([[24198,  3291],
        [ 3171, 24241]], dtype=int64)
```

Decision Tree Classifier

```
In [166... # activating DecisionTree Classifier
model_dtc=DecisionTreeClassifier()

# passing xr_train_scaled, yr_train to training the model
model_dtc.fit(Xr_train_scaled,Yr_train)

model_dtc
```

```
Out[166... ▼ DecisionTreeClassifier ⓘ ?
DecisionTreeClassifier()
```

```
In [167... y_pred_dtc=model_dtc.predict(Xr_test_scaled) # predicting yr_test data
```

```
In [168... # classification report for decisionTreeClassifier

print(classification_report(y_pred_dtc,Yr_test))
```

	precision	recall	f1-score	support
0	0.69	0.99	0.82	19159
1	0.99	0.77	0.86	35742
accuracy			0.84	54901
macro avg	0.84	0.88	0.84	54901
weighted avg	0.89	0.84	0.85	54901

```
In [169... confusion_matrix(y_pred_dtc,Yr_test)
```

```
Out[169... array([[18978,  181],
        [ 8391, 27351]], dtype=int64)
```

RandomForestClassifier()


```
In [170...] model_rfc=RandomForestClassifier() #activating the fuction
model_rfc.fit(Xr_train_scaled,Yr_train)
```

```
Out[170...] RandomForestClassifier
RandomForestClassifier()
```

```
In [171...] y_pred_rfc=model_rfc.predict(Xr_test_scaled)
```

```
In [172...] print(classification_report(y_pred_rfc,Yr_test))
```

	precision	recall	f1-score	support
0	0.86	0.99	0.92	23757
1	0.99	0.87	0.93	31144
accuracy			0.92	54901
macro avg	0.92	0.93	0.92	54901
weighted avg	0.93	0.92	0.92	54901

```
In [173...] confusion_matrix(y_pred_rfc,Yr_test)
```

```
Out[173...] array([[23415, 342],
       [ 3954, 27190]], dtype=int64)
```

XGBOOST

```
In [174...] model_xgb=XGBClassifier()
model_xgb.fit(Xr_train_scaled,Yr_train)
```

```
Out[174...] XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rou
               nds=None,
               enable_categorical=False, eval_metric=None, feature_ty
               pes=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_
               bin=None,
```

```
In [175...] y_pred_xgb=model_xgb.predict(Xr_test_scaled)
```

```
In [176...] print(classification_report(y_pred_xgb,Yr_test))
```

	precision	recall	f1-score	support
0	0.67	0.97	0.79	19008
1	0.98	0.75	0.85	35893
accuracy			0.83	54901
macro avg	0.82	0.86	0.82	54901
weighted avg	0.87	0.83	0.83	54901

In [177... `confusion_matrix(y_pred_xgb,Yr_test)`

Out[177... `array([[18386, 622],
[8983, 26910]], dtype=int64)`

finding the hyperparameter tuning and best param grid

```
In [178... from sklearn.model_selection import GridSearchCV, cross_val_score
from sklearn.linear_model import LogisticRegression

# Define the parameter grid to search over
param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100], # Regularization parameter
    'penalty': ['l1', 'l2'] # Penalty type
}

# Create a Logistic Regression model
logistic = LogisticRegression()

# Create a GridSearchCV object
grid_search = GridSearchCV(estimator=logistic, param_grid=param_grid, cv=10)

# Initialize an empty list to store the accuracy scores
accuracy_scores = []

# Perform cross-validation 10 times
for _ in range(10):
    # Fit the GridSearchCV object to the training data
    grid_search.fit(Xr_train_scaled, Yr_train)

    # Get the best parameters
    best_params = grid_search.best_params_

    # Perform cross-validation with the best model
    cv_scores = cross_val_score(grid_search.best_estimator_, Xr_train_scaled, Yr

    # Store the mean accuracy score
    accuracy_scores.append(cv_scores.mean())

# Print the accuracy scores obtained over 10 iterations
#print("Accuracy scores over 10 iterations:", accuracy_scores)
print("Accuracy scores over 10 iterations:", [{":.2f}".format(score) for score in accuracy_scores])

# Get the best parameters and best score
best_params = grid_search.best_params_
best_score = grid_search.best_score_
```

```
print("Best parameters found:", best_params)
print("Best cross-validation score:", best_score)
```

Accuracy scores over 10 iterations: ['0.89', '0.89', '0.89', '0.89', '0.89', '0.89', '0.89', '0.89', '0.89', '0.89']
 Best parameters found: {'C': 0.001, 'penalty': 'l2'}
 Best cross-validation score: 0.8853698959839237

Final Model

```
In [179... from sklearn.linear_model import LogisticRegression

# Create a Logistic Regression model with the best parameters
final_model = LogisticRegression(C=0.001, penalty='l2')

# Fit the final model to the entire training dataset
final_model.fit(Xr_train_scaled, Yr_train)
```

```
Out[179... LogisticRegression ⓘ ?
LogisticRegression(C=0.001)
```

```
In [180... import pickle

# Save the final model to a pickle file
with open('final_model.pkl', 'wb') as file:
    pickle.dump(final_model, file)
```

```
In [181... import pickle
import numpy as np

# Load the model from the pickle file
with open('final_model.pkl', 'rb') as file:
    loaded_model = pickle.load(file)

# Define the mean and standard deviation of the training data
mean_values = [41.885856, 0.07485, 0.03942, 27.320767, 5.527507, 138.058060]
std_values = [22.516840, 0.26315, 0.194593, 6.636783, 1.070672, 40.708136]

# Define the input features for prediction
age = 30
hypertension = 0
heart_disease = 0
bmi = 100.0
HbA1c_level = 5.0
blood_glucose_level = 90

# Scale the input features manually
scaled_features = [(x - mean) / std for x, mean, std in zip(
    [age, hypertension, heart_disease, bmi, HbA1c_level, blood_glucose_level],
    mean_values, std_values
)]

# Make predictions on the scaled data
prediction = loaded_model.predict([scaled_features])
```

```
# Print the prediction
if prediction[0] == 1:
    print("Diabetic")
else:
    print("Not Diabetic")
```

Diabetic